



# Chunked Historical Processing Flow

This document outlines the architecture, control flow, and technical implementation of chunked historical data processing in EventTrader.

## 1. Overall System Architecture

The chunked historical processing system uses a sequential, multi-stage approach to reliably process large historical date ranges while managing system resources effectively.

### High-Level Flow

1. **Bash Script:** Entry point via `event_trader.sh chunked-historical`
2. **Date Chunking:** Breaks date range into smaller chunks (default: 5 days per chunk)
3. **Sequential Processing:** Launches separate Python processes for each chunk
4. **Completion Monitoring:** Python monitors Redis for processing completion
5. **Clean State:** Ensures clean system termination between chunks
6. **Embedding Generation:** Triggers vector embedding creation after processing
7. **Reconciliation:** Handles edge cases where items remain stuck

```
chunked-historical {from_date} {to_date}
|
|— Chunk 1: Process days 1-5
|   |— Monitor completion via Redis state
|
|— Chunk 2: Process days 6-10
|   |— Monitor completion via Redis state
|
|— Chunk N: Process remaining days
|   |— Monitor completion via Redis state
```

## 2. Entry Point: `event_trader.sh` in Detail

- **`event_trader.sh`:** Bash script control interface
  - `chunked-historical` command
    - **Command validation:**
      - Checks that required `FROM_DATE` is provided
      - Uses default `TO_DATE` (today) if not specified
      - Verifies date format (YYYY-MM-DD)
    - **`process_chunked_historical()` function execution:**
      - **Configuration loading:**
        - Calls `detect_python()` to find Python interpreter

- Loads `HISTORICAL_CHUNK_DAYS` and `HISTORICAL_STABILITY_WAIT_SECONDS` from `config/feature_flags.py`
- Validates configuration values are positive integers
- **Logging setup:**
  - Creates unique folder `logs/ChunkHist_{FROM_DATE}_{TO_DATE}_{TIMESTAMP}/`
  - Creates separate log files:
    - Combined shell log file for tracking overall process
    - Individual log file for each chunk's Python process
  - Defines `shell_log()` function to write to shell log
  - Shell logs are written to combined file while Python logs go to chunk-specific files
- **System checks:**
  - Verifies Redis connectivity with `redis-cli ping`
  - Logs available data sources (news, reports, transcripts)
  - Records total start time for duration tracking
- **Date chunking:**
  - Converts date strings to Unix timestamps (OS-specific compatibility)
  - Creates chunks based on `HISTORICAL_CHUNK_DAYS` configuration
  - Initializes chunk counter and monitoring variables
- **Processing loop** (for each chunk):

- Records chunk start time
- Calculates chunk start/end dates
- Creates chunk-specific log file
- Executes `stop-all` to terminate previous instances
- Executes Python processor with parameters:

```
$PYTHON_CMD "$SCRIPT_PATH" \
--from-date "$chunk_start" \
--to-date "$chunk_end" \
-historical \
--ensure-neo4j-initialized \
--log-file "$CHUNK_LOG_FILE"
```

- **Process monitoring and PID tracking:**
  - Captures and stores Python PID: `EVENTTRADER_PID=$!`
  - Writes PID to file for external tracking
  - Monitors process with timeout controls:
    - Watches for completion messages in log files
    - Periodically checks if process is still running via PID
    - Times out after 2 hours maximum per chunk
  - Handles process termination if needed:
    - First attempts graceful shutdown with `SIGTERM`
    - Waits 5 seconds for clean exit
    - Forces termination with `SIGKILL` if process remains

- Captures Python exit code
- Handles success/failure scenarios
- **Log summary extraction:**
  - Extracts key events (errors, warnings, completions) from chunk log
  - Appends summary to combined log file
- Executes `stop-all` to ensure clean state between chunks
- Calculates and logs chunk duration
- Advances to next chunk start date
- **Finalization:**
  - Calculates and logs total process duration
  - Creates summary file with statistics
  - Logs completion message with full range processed

### 3. Python Application: `run_event_trader.py`

- **`run_event_trader.py`:** Main Python entry point
  - **`main()` function:**
    - **Error handling setup:**
      - Comprehensive try-except block for entire application
      - Graceful shutdown on errors with detailed logging
    - **Command-line processing:**
      - `parse_args()`: Processes command-line arguments
        - Parses `--from-date` and `--to-date` (required)
        - Handles `-historical` flag to disable live data
        - Processes `--ensure-neo4j-initialized` flag
        - Configures `--log-file` path (receives chunk-specific log file)
    - **Feature flag configuration:**
      - Sets `ENABLE_HISTORICAL_DATA=True` (for chunked-historical mode)
      - Sets `ENABLE_LIVE_DATA=False` (historical only)
    - **Logging initialization:**
      - Sets up logging framework with file and console handlers
      - Writes to the chunk-specific log file provided by shell script
    - **Signal handlers:**
      - Registers handlers for SIGINT/SIGTERM for clean shutdown
    - **DataManager creation:**
      - Initializes manager with date range: `manager = DataManager(date_from, date_to)`
    - **Neo4j validation:**
      - Verifies Neo4j connection with `manager.has_neo4j()`
      - Proceeds if initialized, exits with error if failed
    - **System startup:**
      - Calls `manager.start()` to begin processing
      - Enters monitoring loop for completion in historical-only mode

- **Completion monitoring** (historical mode):
  - Helper functions:
    - `check_initial_processing()`: Checks fetch completion and queue states
    - `only_withreturns_remain()`: Detects special case where only items in withreturns remain
  - Obtains Redis connection from manager
  - Monitors multiple Redis indicators for each source:
    1. Batch fetch completion flags: `batch:{source}:{from}-{to}:fetch_complete`
    2. Raw queue emptiness: `{source}:queues:raw`
    3. Historical namespace emptiness: `{source}:hist:{raw|processed}:*`
    4. Pending returns sets: `{source}:pending_returns`
    5. WithReturns namespace: `{source}:withreturns:*`
    6. WithoutReturns namespace: `{source}:withoutreturns:*`
  - Periodically checks status every 30 seconds
  - Implements timeout/retry mechanism (`WITHRETURNS_MAX_RETRIES`)
  - Triggers reconciliation if processing stalls with only withreturns items
  - Logs completion status for each source
- **Embedding generation:**
  - After all processing completes, calls `neo4j_processor.batch_process_qaexchange_embeddings()`
  - Uses batch size defined in `feature_flags.QAEXCHANGE_EMBEDDING_BATCH_SIZE`
  - Embeds QA pairs to enable semantic search
- **Shutdown and cleanup:**
  - Calls `manager.stop()` to shut down cleanly
  - Logs “Shutdown complete. Exiting Python process” (triggers shell script detection)
  - Exits Python process with success code (0)
  - Returns control to bash script for next chunk

## 4. Data Manager Initialization and Source Management

- **\*\*DataManager.\_\_init\_\_(date\_from, date\_to)\*\*:**
  - **Core initialization:**
    - Sets up logging and signal handlers
    - Stores date range in `historical_range` dictionary
    - Creates empty `sources` dictionary
  - **Source initialization:**
    - `initialize_sources()`: Creates source manager instances
      - `BenzingaNewsManager(historical_range)`: News data source
        - Initializes Redis connections (separate for live/historical)
        - Creates NewsProcessor for raw→processed conversion
        - Creates ReturnsProcessor for market impact calculation
      - `ReportsManager(historical_range)`: SEC filings source
        - Initializes Redis connections and processors

- `TranscriptsManager(historical_range):` Earnings calls source
  - Initializes Redis connections and processors
- **Neo4j initialization:**
  - `initialize_neo4j():` Sets up graph database connection
    - Creates Neo4jProcessor instance
    - Initializes database schema (creates constraints, indexes)
    - Creates date nodes for the processing date range
    - Validates connectivity
    - Starts background processing thread via `process_with_pubsub()`

## 5. Source Manager Start Processes

- **manager.start():** Initiates all data processing
  - Calls `start()` on each source manager, which:
    - **BenzingaNewsManager.start():**
      - `rest_client.get_historical_data():` Fetches data directly in the main execution path (not in a separate thread)
      - Creates and starts daemon threads:
        - `processor_thread:` Runs `processor.process_all_news()`
        - `returns_thread:` Runs `returns_processor.process_all_returns()`
    - **ReportsManager.start():**
      - Creates and starts daemon threads:
        - `processor_thread:` Runs `processor.process_all_reports()`
        - `returns_thread:` Runs `returns_processor.process_all_returns()`
        - `historical_thread:` Runs `rest_client.get_historical_data()` (unlike News, Reports data is fetched in a thread)
    - **TranscriptsManager.start():**
      - `_initialize_transcript_schedule():` Sets up retrieval plan
      - Creates and starts daemon threads:
        - `processor_thread:` Runs `processor.process_all_transcripts()`
        - `returns_thread:` Runs `returns_processor.process_all_returns()`
        - `historical_thread:` Runs `_fetch_historical_data()`
  - Returns dictionary of status results from all source starts

## 6. Processor Data Flow

### 6.1 Per-Source Processing Flow

Each data source follows a similar processing pattern:

1. **Fetch Stage** - Historical data retrieval
  - Sets `batch:{source}:{from}-{to}:fetch_complete` when fetch completes

- Stores raw items in Redis: `{source}:hist:raw:{id}`, adds to `{source}:queues:raw`
- 2. Base Processing - Raw to structured data conversion**
  - `BaseProcessor.process_all_items()` consumes items from `{source}:queues:raw`
  - Converts raw data to structured format (standardize, clean, add metadata)
  - Stores processed items: `{source}:hist:processed:{id}`
  - Adds to `{source}:queues:processed`
- 3. Returns Processing - Market impact analysis**
  - `ReturnsProcessor.process_all_returns()` handles processed items
  - Calculates hourly, session, and daily returns
  - Uses event metadata to determine return timing
  - Stores as `{source}:withreturns:{id}` (complete) or `{source}:withoutreturns:{id}` (pending)
  - Adds incomplete items to `{source}:pending_returns` ZSET with due timestamp
- 4. Neo4j Integration - Graph database storage**
  - `Neo4jProcessor.process_with_pubsub()` processes `withreturns/withoutreturns` entries
  - Creates nodes and relationships in Neo4j database
  - Deletes items from Redis once successfully stored in Neo4j

## 6.2 Embedding Generation

After all processing completes, the system triggers embedding generation:

- **Embedding Generation Process:**
  - `batch_process_qaexchange_embeddings()`: Processes all question-answer pairs
  - Creates vector embeddings using OpenAI API
  - Stores embeddings as node properties in Neo4j
  - Enables semantic similarity search on text content

## 7. Completion Monitoring and Error Handling

### 7.1 Completion Monitoring

The system employs a multi-stage approach to determine when processing is complete:

- **Fetch Completion:**
  - Checks `batch:{source}:{from}-{to}:fetch_complete` flags
  - Verifies `{source}:queues:raw` is empty
  - Ensures no items remain in historical raw namespace
- **Processing Completion:**
  - Verifies historical processed namespace is empty
  - Checks pending returns sets are empty
- **Final Completion:**
  - Ensures `withreturns` and `withoutreturns` namespaces are empty
  - Indicates all items have been successfully moved to Neo4j

## 7.2 Error Handling and Recovery

The system implements several error recovery mechanisms:

- **Timeout Controls:**
  - Maximum 2-hour timeout per chunk in shell script
  - Monitoring cycle with retries in Python process
- **Reconciliation:**
  - After `WITHRETURNS_MAX_RETRIES` cycles (default: 3) monitoring cycles
  - Checks if only `withreturns` items remain (`only_withreturns_remain()`)
  - Triggers `reconcile_missing_items()` to force reload from Redis to Neo4j
- **Clean Shutdown:**
  - Signal handling for graceful termination
  - `manager.stop()` ensures proper cleanup
  - `stop-all` command between chunks

## 8. Thread Execution By Mode

The following table shows which threads are started (✅) or not started (❌) in chunked-historical mode compared to live mode:

Thread	Live Mode (-live)	Chunked-Historical (-historical)
processor_thread (News)	✅	✅
returns_thread (News)	✅	✅
ws_thread (News WebSocket)	✅	❌
historical_thread (News Historical Fetch)	❌	❌
processor_thread (Reports)	✅	✅
returns_thread (Reports)	✅	✅
ws_thread (Reports WebSocket)	✅	❌
historical_thread (Reports Historical Fetch)	❌	✅
processor_thread (Transcripts)	✅	✅
returns_thread (Transcripts)	✅	✅
ws_thread (Transcripts)	❌	❌
historical_thread (Transcripts Historical Fetch)	❌	✅
neo4j_thread (PubSub Event Processor)	✅	✅

## 9. Detailed Redis Flow and Completion States

### 9.1 Fetch Completion Indicators

For each source, the system sets completion flags when fetching is complete:

```
batch:news:{from_date}-{to_date}:fetch_complete = "1"  
batch:reports:{from_date}-{to_date}:fetch_complete = "1"  
batch:transcripts:{from_date}-{to_date}:fetch_complete = "1"
```

### 9.2 Queue and Namespace Emptiness Checks

The system verifies multiple Redis structures for emptiness:

1. **Raw Queues** - Must be empty to indicate all raw data has been processed:

```
news:queues:raw  
reports:queues:raw  
transcripts:queues:raw
```

2. **Historical Namespaces** - Must be empty to ensure all items were processed:

```
news:hist:raw:*  
news:hist:processed:*  
reports:hist:raw:*  
reports:hist:processed:*  
transcripts:hist:raw:*  
transcripts:hist:processed:*
```

3. **Return Storage** - All items must be processed by Neo4j:

```
news:withreturns:*  
news:withoutreturns:*  
reports:withreturns:*  
reports:withoutreturns:*  
transcripts:withreturns:*  
transcripts:withoutreturns:*
```

4. **Pending Returns** - ZSET must be empty to indicate all returns are calculated:

```
news:pending_returns  
reports:pending_returns
```

### 9.3 Reconciliation Logic

When only withreturns items remain:

1. System detects withreturns keys but all other conditions are met
2. After WITHRETURNS\_MAX\_RETRIES cycles (default: 3)
3. Triggers reconcile\_missing\_items()
4. Forces check and reload of Redis keys into Neo4j
5. Handles potential race conditions if Neo4j connection issues occurred



## 10. Performance and Scalability Considerations

### 10.1 Resource Management

- **Memory Efficiency:**
  - Processing in chunks prevents memory exhaustion
  - Clean termination between chunks releases memory
- **CPU Utilization:**
  - Limit of one chunk processed at a time
  - Parallel processing within each chunk via daemon threads
- **API Rate Limiting:**
  - Smaller chunks reduce burst API usage
  - Configurable chunk size via `HISTORICAL_CHUNK_DAYS`

### 10.2 Failure Isolation

- **Chunk Isolation:**
  - Failure in one chunk doesn't affect others
  - Shell script tracks per-chunk success/failure
  - Detailed logs for troubleshooting each chunk

### 10.3 Configuration Parameters

- **`HISTORICAL_CHUNK_DAYS`** (default: 5):
  - Controls chunk size in days
  - Lower values reduce memory usage but increase overhead
  - Higher values increase efficiency but require more resources
- **`HISTORICAL_STABILITY_WAIT_SECONDS`** (default: 60):
  - Optional wait time between chunks
  - Allows system stability before starting new chunk
- **`WITHRETURNS_MAX_RETRIES`** (default: 3):
  - Controls reconciliation trigger threshold
  - Number of monitoring cycles before forcing reconciliation

## 11. Logging and Monitoring

### 11.1 Log Structure

- **Shell Script Logs:**
  - Main shell operations log: `logs/ChunkHist_{FROM_DATE}_to_{TO_DATE}_{TIMESTAMP}/combined_{FROM_DATE}_to_{TO_DATE}.log`
    - Contains shell script operations and commands
    - Includes summaries extracted from chunk logs
    - Tracks overall progress across all chunks
  - Per-chunk logs: `logs/ChunkHist_{FROM_DATE}_to_{TO_DATE}_{TIMESTAMP}/chunk_{start}_to_{end}.log`
    - Contains detailed Python process logs for each chunk
    - Detailed error messages and stack traces

- Full Redis completion monitoring status
- Summary file: `logs/ChunkHist_{FROM_DATE}_to_{TO_DATE}_{TIMESTAMP}/summary.txt`
  - Provides overview of complete run
- **Log Flow:**
  - Shell script writes directly to combined log via `shell_log()`
  - Python processes write to their individual chunk logs
  - Shell script extracts important events from chunk logs and appends to combined log
  - This two-tier approach keeps detailed logs separate while maintaining overall visibility

## 11.2 Monitoring Points

- **Critical Checkpoints:**
  1. Fetch completion for each source
  2. Queue emptiness
  3. Returns calculation completion
  4. Neo4j ingestion completion
  5. Embedding generation success
- **Performance Metrics:**
  - Per-chunk processing time
  - Overall job duration
  - Items processed per source

## 12. Conclusion

The chunked historical processing system provides a robust, scalable approach to processing large historical data ranges. Its key advantages include:

1. **Resource Efficiency:** Controlled memory and CPU usage
2. **Fault Tolerance:** Chunk isolation prevents cascading failures
3. **Comprehensive Monitoring:** Multi-stage completion checks
4. **Error Recovery:** Automatic reconciliation mechanisms
5. **Scalability:** Configurable parameters for different environments

These capabilities enable reliable processing of extensive historical data while maintaining system stability and performance.