

Comprehensive Redis Flow Description — Live vs Historical

This document outlines how Redis keys and queues are created, moved, scheduled, or deleted across the **News**, **SEC Reports**, and **Transcripts** pipelines. Each section clearly separates the **Live** and **Historical** workflows, linking Redis actions to their respective scripts and logic.

1. News (Benzinga)

a) WebSocket (**bz_websocket.py**) — Live

- Connects to Benzinga WebSocket stream.
- For each real-time news item:
 - Uses `set_news()`:
 - * Stores raw JSON:
 - `news:benzinga:live:raw:{id}.{updated}`
 - * Pushes key to:
 - `news:benzinga:queues:raw`
 - * Publishes to Redis (optional):
 - `news:benzinga:raw:channel`

b) REST API (**bz_restAPI.py**) — Historical

- Fetches historical news using batch REST queries.
 - Uses `set_news_batch()`:
 - Stores each news item as:
 - * `news:hist:raw:{id}.{updated}`
 - Pushes key to:
 - * `news:queues:raw`
 - On success:
 - * Sets: `batch:news:{from}-{to}:fetch_complete`
-

News Processing (**NewsProcessor.py**) — Live + Historical

- Inherits `BaseProcessor`, monitors:
 - `news:queues:raw` (both live and hist share this queue)
- Each raw key is:
 1. Fetched from Redis
 2. Validated for allowed symbols
 3. Cleaned and normalized
 4. Augmented using `EventManager.process_event_metadata()`
 5. Stored to:
 - `news:benzinga:{live|hist}:processed:{id}.{updated}`
 6. Pushed to:
 - `news:benzinga:queues:processed`
 7. If live:
 - Publishes to: `news:benzinga:live:processed`

Returns Processing (**ReturnsProcessor.py**) — Live + Historical

- For **Live**: Subscribes to:
 - `news:benzinga:live:processed`
 - For **Historical**: Scans and processes in batches:
 - `news:hist:processed:*`
 - For each processed item:
 1. Calculates:
 - Hourly, Session, and Daily returns using `market_session.py`
 2. Annotates `metadata.returns_schedule` and actual return values
 3. Writes to:
 - `news:withreturns:{id}` (*if complete*) or
 - `news:withoutreturns:{id}` (*if pending*)
 4. If incomplete:
 - Adds to `news:pending_returns` (ZSET)
 - Periodic checks pull from the ZSET:
 - When returns become available, `withoutreturns` → `withreturns` + delete original
-

2. SEC Reports

a) WebSocket (**sec_websocket.py**) — Live

- Receives live 8-K, 10-Q, 10-K filings.
- Uses `set_filing()`:
 - Writes:
 - * `reports:live:raw:{accessionNo}.{filedAt}`
 - Pushes to:
 - * `reports:queues:raw`

b) REST API (**sec_restAPI.py**) — Historical

- Fetches historical filings.
 - Uses `set_filing()` inside batch logic:
 - Stores under:
 - * `reports:hist:raw:{accessionNo}.{filedAt}`
 - Pushes to:
 - * `reports:queues:raw`
 - Completion flag:
 - * `batch:reports:{from}-{to}:fetch_complete`
-

Report Processing (**ReportProcessor.py**) — Live + Historical

- Inherits `BaseProcessor`, listens to:
 - `reports:queues:raw`

- Each report:
 1. Validated and converted to `UnifiedReport`
 2. Enriched using `EventReturnsManager.process_event_metadata()`
 3. Stored under:
 - `reports:{live|hist}:processed:{accessionNo}.{filedAt}`
 4. Pushed to:
 - `reports:queues:processed`
 5. If live:
 - Publishes to: `reports:live:processed`
-

Returns Processing (**ReturnsProcessor.py**) — Live + Historical

- For **Live**: Subscribes to:
 - `reports:live:processed`
 - For **Historical**: Scans and processes in batches:
 - `reports:hist:processed:*`
 - Same logic as news:
 - Updates `reports:withreturns` or `reports:withoutreturns`
 - Adds to:
 - * `reports:pending_returns` (ZSET)
-

3. Transcripts

a) REST API (**EarningsCallTranscripts.py**) — Live & Historical

- Fetches and standardizes transcripts to `UnifiedTranscript`
 - Calls `store_transcript_in_redis()`:
 - Sets:
 - * `transcripts:{live|hist}:raw:{symbol}_{timestamp}`
 - Publishes to:
 - * `admin:transcripts:notifications`
 - Also:
 - * Pushed to standard `transcripts:queues:raw` for `BaseProcessor` compatibility
-

Transcript Processing (**TranscriptProcessor.py**) — Live + Historical

- **Combined Approach:**
 - Uses **both** ZSET scheduling and queue processing
 - Runs a scheduling thread for ZSET monitoring in parallel with `BaseProcessor`
 - Controlled by `ENABLE_LIVE_DATA` feature flag
- **Scheduling Thread:**
 - Monitors `admin:transcripts:schedule` ZSET with scores as timestamps
 - Processes up to 5 due transcripts at once
 - Performs daily date transition checks
 - Sleep timeout limited by `MAX_TRANSCRIPT_SLEEP_SECONDS` (300 seconds)

- Listens for notifications on `admin:transcripts:notifications` channel
 - For each due item in ZSET:
 1. Checks if already processed via `admin:transcripts:processed` SET
 2. Verifies symbol is in tradable universe
 3. Fetches transcript via `EarningsCallProcessor` for exact date
 4. Falls back to current date if no transcript found
 5. Validates fields and standardizes format for `BaseProcessor`
 6. Enriches with `EventReturnsManager.process_event_metadata()`
 7. Stores in both:
 - `transcripts:{live|hist}:processed:{symbol}_{timestamp}`
 - `transcripts:withreturns:{id}` or `transcripts:withoutreturns:{id}`
 - **State Management:**
 - Removes from `admin:transcripts:schedule` ZSET when completed (via `zrem`)
 - Adds to `admin:transcripts:processed` SET for deduplication (via `sadd`)
 - Reschedules via ZSET if not ready yet (5-minute intervals) (via `zadd`)
 - Has error recovery with 30-minute retry on failures
 - Cleans up raw keys after successful processing
 - Publishes notification messages to:
 - * `admin:transcripts:notifications` with `processed:{event_key}` or `rescheduled:{event_key}` prefixes
-

Reconciliation Logic (**reconcile.py**)

- Periodic or triggered reconciliation of missed items.
 - Scans:
 - `news:withreturns:*`
 - `reports:withreturns:*`
 - `transcripts:withreturns:*`
 - Verifies if node exists in Neo4j:
 - If yes → deletes from Redis
 - If no → re-pushes to `queues:processed`
 - Fallbacks:
 - `metadata.created` + symbol used to match Redis and DB state
 - Optional TTL cleanup logic may remove items after stale duration.
-

Event Returns (**EventReturnsManager.py**)

- Used in all 3 pipelines.
- Computes:
 - `returns_schedule = {hourly, session, daily}`
 - Each with start/end timestamps
- Logic adapts to:
 - Post-market vs pre-market vs in-session events
- Returns stored in metadata:
 - Under `event`, `returns_schedule`, and `{hourly|session|daily}_return`

Redis Structures Summary

Queues:

- `news:queues:raw`
- `news:benzinga:queues:raw`
- `news:benzinga:queues:processed`
- `reports:queues:raw`
- `reports:queues:processed`
- `transcripts:queues:raw` (used for BaseProcessor compatibility)
- `transcripts:queues:processed` (used for BaseProcessor compatibility)

Pub/Sub Channels:

- `news:benzinga:live:processed`
- `reports:live:processed`
- `admin:transcripts:notifications`

ZSETs:

- `news:pending_returns`
- `reports:pending_returns`
- `admin:transcripts:schedule`

Key Patterns:

Raw

- `news:benzinga:live:raw:{id}.{updated}`
- `news:hist:raw:{id}.{updated}`
- `reports:live:raw:{accessionNo}.{filedAt}`
- `reports:hist:raw:{accessionNo}.{filedAt}`
- `transcripts:live:raw:{symbol}_{timestamp}`
- `transcripts:hist:raw:{symbol}_{timestamp}`

Processed

- `news:benzinga:live:processed:{id}.{updated}`
- `news:hist:processed:{id}.{updated}`
- `reports:live:processed:{accessionNo}.{filedAt}`
- `reports:hist:processed:{accessionNo}.{filedAt}`
- `transcripts:live:processed:{symbol}_{timestamp}`
- `transcripts:hist:processed:{symbol}_{timestamp}`

Final

- `news:withreturns:{id} / news:withoutreturns:{id}`
- `reports:withreturns:{id} / reports:withoutreturns:{id}`
- `transcripts:withreturns:{id} / transcripts:withoutreturns:{id}`

Set Membership

- `admin:transcripts:processed`

Admin

- `admin:tradable_universe:stock_universe`
 - `admin:tradable_universe:symbols`
 - `admin:news:last_message_time`
 - `admin:reports:last_message_time`
 - `admin:news:shutdown_state`
 - `admin:reports:shutdown_state`
 - `admin:backfill:news_restart_gap`
 - `admin:backfill:reports_restart_gap`
-